
Shaoqi Wang

I am a fourth-year Ph.D. student in University of Colorado, Colorado Springs. My research interests include Big Data Analytics, Distributed ML and DL Systems. Currently, I am working on building dynamic resource scheduler for multiple ML/DL training jobs on clusters.

Email: swang@uccs.edu Phone: (+1)7193325477

<https://ibingoogle.github.io/swang/index.html>

<https://www.linkedin.com/in/shaoqi-wang-8b392b105/>



Education:

- Pursuing Ph.D. in Computer Science
University of Colorado, Colorado Springs, USA, Fall 2015 - present
- Master Degree in Engineering
University of Science and Technology of China, China, Fall 2012 - Spring 2015
- Bachelor Degree in Computer Science
Anhui Normal University, China, Fall 2008 - Spring 2012

Project Experiences:

- Research Assistant
DISCO Lab, University of Colorado, Colorado Springs, USA, Fall 2015 – present
 1. **Dynamic Resource Scheduler:** Dec 2018 – Present
The hyper-parameter search for ML/DL model involves multiple ML/DL jobs. Existing cluster schedulers with static resource scheduling are largely not tailored to multiple ML/DL jobs. To this end, currently, I am working on building dynamic resource scheduler for multiple ML/DL jobs on clusters. I plan to implement the scheduler in Tencent Angel and Google TensorFlow.
 2. **iBatch:** Jun 2018 – Present
Scalability of distributed DL training with parameter server architecture is often communication constrained in large clusters. To this end, I proposed iBatch, a novel communication approach that batches parameter communication and forward computation to overlap them with each other. I implemented iBatch in the distributed DL system Intel BigDL. This work is to appear in AAAI 2019.
 3. **A-BSP:** Aug 2017 – Aug 2018
Executing distributed ML jobs on Apache Spark follows BSP model, in which the synchronization is significantly delayed by stragglers. To this end, I proposed a novel BSP-based Aggressive synchronization (A-BSP) model based on the convergent property of iterative ML algorithms, by allowing the algorithm to use the updates generated based on partial input data for synchronization. I implemented A-BSP as a light-weight BSP-compatible mechanism in Apache Spark and also extended A-BSP onto Petuum system. This work is published in ACM/IFIP Middleware 2018.
 4. **FlexPara:** Feb 2017 – Dec 2017
Computational skewness is a significant challenge in multi-tenant data-parallel clusters that introduce dynamic heterogeneity of machine capacity in distributed data processing. To this end, I proposed FlexPara, a parameter partition approach that leverages the non-linear relationship and provisions adaptive tasks to match the distinct machine capacity so as to address the skewness in iterative ML jobs on data-parallel clusters. I implemented FlexPara in Apache Spark. This work is to appear in IEEE INFOCOM 2019.
 5. **Dawn:** May 2016 – May 2017
The performance of parallel jobs is often constrained by the cluster's hard-to-scale network bisection bandwidth. To this end, I proposed Dawn, a dependency-aware network-adaptive

scheduler that aggregates and co-locates the data and tasks of dependent jobs to improve data locality. I implemented Dawn on Apache Hadoop and Apache Tez. This work is published in IEEE ICAC 2017 and IEEE TPDS.

- Research Assistant

Institute of Intelligent Machines, Chinese Academy of China, Fall 2013 – Spring 2015

1. **Bilingual Lexicon Extraction:**

I proposed a novel method to improve the bilingual lexicon extraction performance from comparable corpora via optimizing translation candidate lists.

Technical Skills:

- Programming Language: Java, Python, Scala, Shell script
- Big Data Systems: Apache Hadoop, Apache Spark
- Distributed ML/DL Systems: Spark MLlib, Intel BigDL, Tencent Angel, TensorFlow

Languages:

- Chinese: Native
- English: Professional working proficiency

Publications and Technical Reports:

1. Addressing Skewness in Iterative ML Jobs with Parameter Partition
Shaoqi Wang, Wei Chen, Xiaobo Zhou, Sang-Yoon Chang, Mike Ji
in IEEE International Conference on Computer Communications (INFOCOM 2019), Paris, France.
2. Scalable Distributed DL Training: Batching Communication and Computation
Shaoqi Wang, Aidi Pi, Xiaobo Zhou
in AAAI Conference on Artificial Intelligence (AAAI 2019), Honolulu, Hawaii, USA.
3. Aggressive Synchronization with Partial Processing for Iterative ML Jobs on Clusters
Shaoqi Wang, Wei Chen, Aidi Pi, Xiaobo Zhou
in ACM/IFIP International Middleware Conference (Middleware 2018), Rennes, France.
4. Dependency-aware Network Adaptive Scheduling of Data-Intensive Parallel Jobs
Shaoqi Wang, Wei Chen, Xiaobo Zhou, Liqiang Zhang, Yin Wang
in IEEE Transactions on Parallel and Distributed Systems (TPDS), August 2018.
4. Scalable Distributed Machine Learning on Data-Parallel Clusters
Shaoqi Wang
Thesis Proposal, Jan 2018.
6. Performance Isolation of Data-Intensive Scale-out Applications in a Multi-tenant Cloud
Palden Lama, **Shaoqi Wang**, Xiaobo Zhou, Dazhao Cheng
in IEEE International Parallel and Distributed Processing Symposium (IPDPS 2018), Canada.
7. Characterizing Scheduling Delay for Low-latency Data Analytics Workloads
Wei Chen, Aidi Pi, **Shaoqi Wang**, Xiaobo Zhou
in IEEE International Parallel and Distributed Processing Symposium (IPDPS 2018), Canada.
8. Improving Utilization and Parallelism of Hadoop Cluster by Elastic Containers
Yinggen Xu, Wei Chen, **Shaoqi Wang**, Xiaobo Zhou, Changjun Jiang
in IEEE International Conference on Computer Communications (INFOCOM 2018), Hawaii, USA.
9. Network-Adaptive Scheduling of Data-Intensive Parallel Jobs with Dependencies in Clusters
Shaoqi Wang, Xiaobo Zhou, Liqiang Zhang, Changjun Jiang
in IEEE International Conference on Autonomic Computing (ICAC 2017), Columbus, OH, USA